I. Two attention pathways carry motion-token content into the action tokens

Swap tasks (ButtonUnmaskSwap | VideoUnmaskSwap), motion model, seed 42, 4 inferences up to first gripper closure; numbers are Button | Video. Layer order: L0–1 language reads memory → by ≈L3 frames hold most of their final motion content → L14–16 action reads frames.

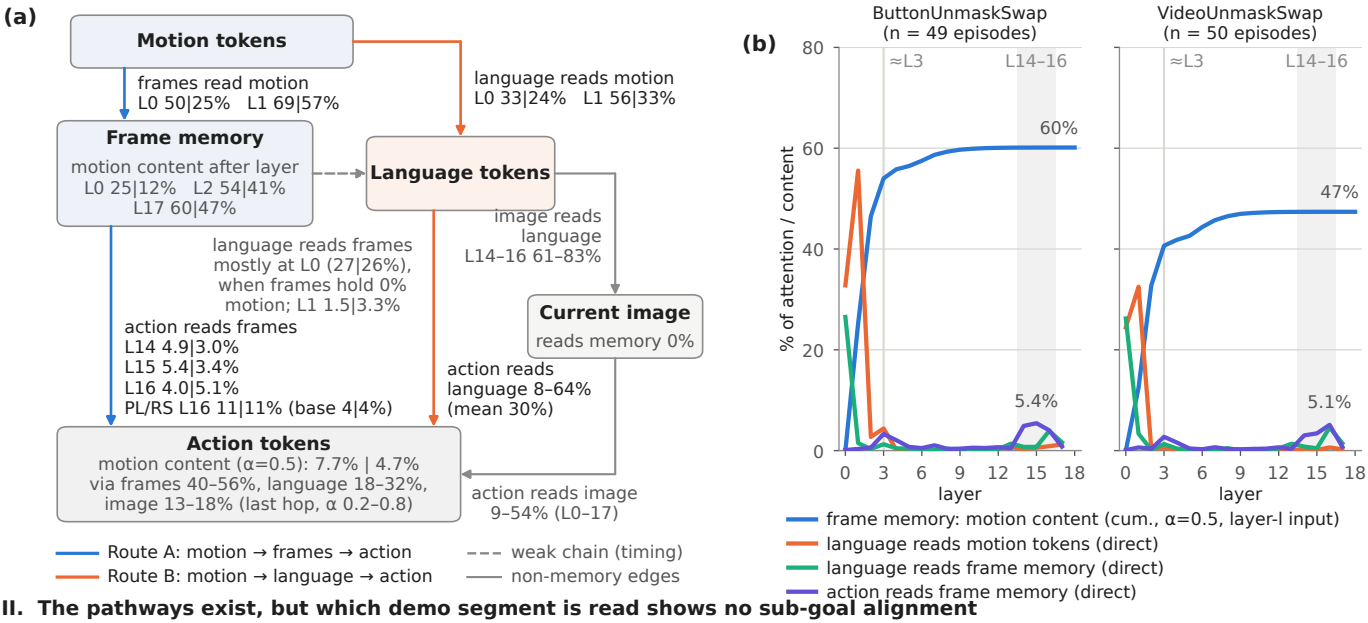

II. The pathways exist, but which demo segment is read shows no sub-goal alignment
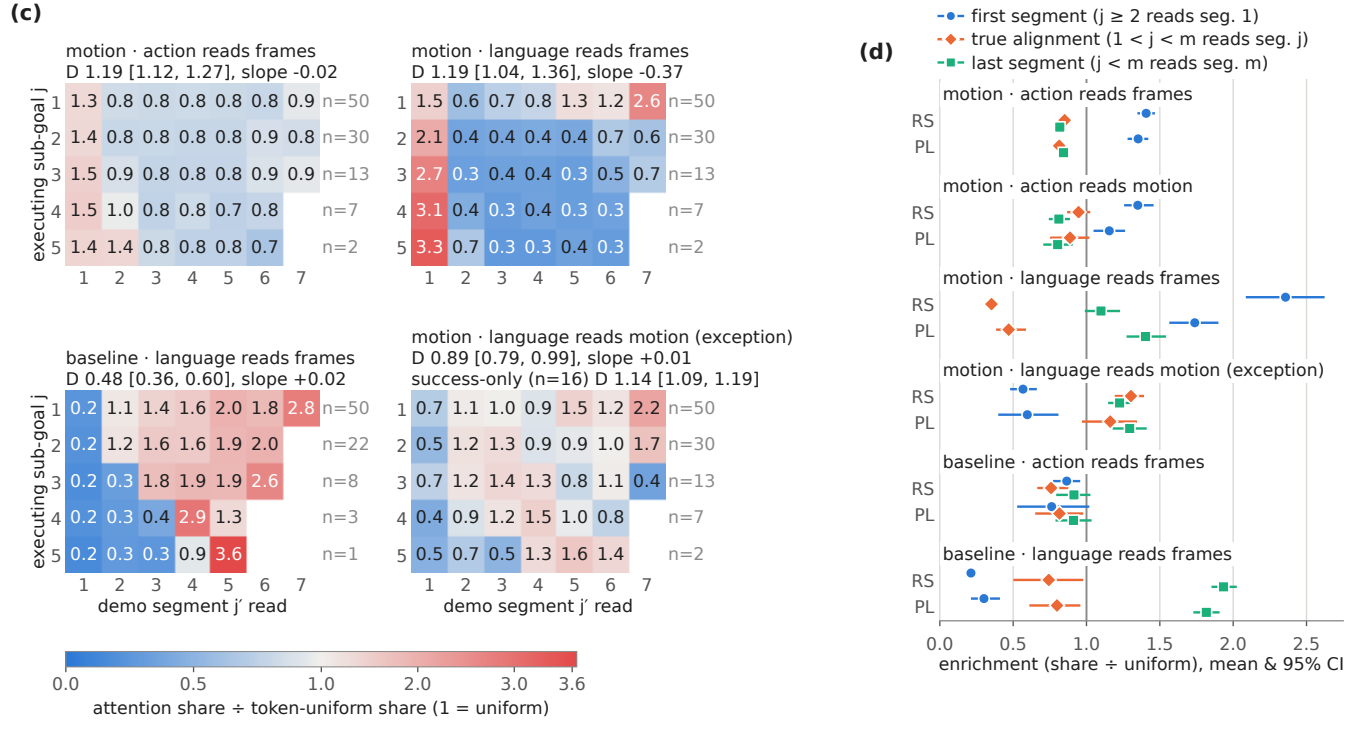
PatternLock / RouteStick, all 922 inferences, seed 42. (c) RouteStick, 50 episodes per model; rows pool episodes with different numbers of sub-goals m (n = episodes per row); frame stream weighted 'single', motion stream 'overlap'; D = diagonal enrichment with episode-level 95% CI. (d) n varies by series (see caption); RS = RouteStick, PL = PatternLock; post-hoc decomposition of D, not corrected for multiple comparisons.